

demo_cross_validation

January 19, 2022

1 Démo : Validation croisée

2 Importer les données

Les données sont disponibles dans le fichier `breast_cancer.csv` dans le répertoire `data`. L'import de données se fait avec la commande `read_csv` de la librairie **pandas**. L'attribut `shape` contient les dimensions de la matrice de données (dataframe).

```
[1]: import pandas as pd

data = pd.read_csv('../data/breast_cancer.csv', sep=';', index_col='id_sample')
print('data', data.shape)
```

data (1016, 51)

La méthode `head()` permet d'afficher les 5 premières lignes du dataframe.

```
[2]: data.head()
```

```
[2]:
```

	NAT1	BIRC5	BAG1	BCL2	BLVRA	CCNB1	\
id_sample							
TCGA-3C-AAAU-01A	7.100449	3.361004	3.972581	4.145669	4.765233	4.788987	
TCGA-3C-AALI-01A	3.453640	4.501040	2.720929	1.493020	5.823480	5.281003	
TCGA-3C-AALJ-01A	4.455574	4.164643	3.911511	4.191457	5.987255	5.229446	
TCGA-3C-AALK-01A	4.297961	3.920234	3.688335	3.894904	5.211594	4.014641	
TCGA-4H-AAAK-01A	1.695378	2.950846	4.110014	3.572843	4.317856	3.772768	

	CCNE1	CDC6	CDC20	CDH3	...	GPR160	\
id_sample					...		
TCGA-3C-AAAU-01A	2.164814	2.633598	4.131205	0.133455	...	4.150233	
TCGA-3C-AALI-01A	2.535437	2.734157	4.176553	0.110023	...	5.561226	
TCGA-3C-AALJ-01A	2.267963	3.379961	4.592752	0.236786	...	2.859309	
TCGA-3C-AALK-01A	0.951107	1.472950	3.806552	0.062392	...	3.063807	
TCGA-4H-AAAK-01A	1.103958	2.338953	3.473484	0.098773	...	3.289418	

	UBE2T	CXXC5	ANLN	CEP55	ACTR3B	MLPH	\
id_sample							
TCGA-3C-AAAU-01A	4.106918	5.528618	3.073409	2.669860	1.928460	5.567999	
TCGA-3C-AALI-01A	5.648057	4.711309	3.881110	3.357553	1.168684	7.064176	

TCGA-3C-AALJ-01A	5.213461	6.152875	2.697093	2.599436	1.177678	5.222420
TCGA-3C-AALK-01A	4.166154	5.612184	2.645664	2.448027	1.026535	6.225590
TCGA-4H-AAAK-01A	3.437585	4.299617	2.068516	2.152652	1.513181	5.485277

	NUF2	TMEM45B	pam50
id_sample			
TCGA-3C-AAAU-01A	2.536764	0.213597	luminal-A
TCGA-3C-AALI-01A	3.124620	3.946538	HER2-enriched
TCGA-3C-AALJ-01A	3.053335	0.281303	luminal-B
TCGA-3C-AALK-01A	1.717959	3.289543	luminal-A
TCGA-4H-AAAK-01A	1.537125	2.976903	luminal-A

[5 rows x 51 columns]

```
[3]: data.groupby(['pam50']).size()
```

```
[3]: pam50
HER2-enriched      82
basal-like         190
luminal-A          543
luminal-B          201
dtype: int64
```

3 Séparer les données d'expression et les étiquettes

```
[4]: # Données d'expression de 50 gènes
X = data.select_dtypes('number')
print('X', X.shape)
```

X (1016, 50)

```
[5]: # Etiquettes correspondantes (sous-types moléculaires)
y = data['pam50']
print('y', y.shape)
```

y (1016,)

4 Créer une validation croisée stratifiée

4.1 Principe

```
[6]: from sklearn.model_selection import StratifiedKFold

random_state = 0
cross_validation = StratifiedKFold(n_splits=3, random_state=random_state,
    ↪ shuffle=True)
```

```

for train_index, test_index in cross_validation.split(X, y):
    X_train = X.iloc[train_index]
    X_test = X.iloc[test_index]
    print('Train', X_train.shape, 'Test', X_test.shape)

```

Train (677, 50) Test (339, 50)

Train (677, 50) Test (339, 50)

Train (678, 50) Test (338, 50)

4.2 Calcul détaillé complet

```

[7]: from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn import metrics

scaler = StandardScaler()
classifier = SVC(kernel='linear', random_state=random_state,
    ↪class_weight='balanced')

accuracy = pd.Series(dtype=float)

iteration = 0
for train_index, test_index in cross_validation.split(X, y):

    iteration += 1 # short version of "iteration = iteration + 1"

    # Train dataset
    X_train = X.iloc[train_index]
    y_train = y.iloc[train_index]

    # Test dataset
    X_test = X.iloc[test_index]
    y_test = y.iloc[test_index]

    # Scaled data
    X_train_scaled = scaler.fit_transform(X_train)
    X_test_scaled = scaler.transform(X_test) # transform only!

    # Train ML classifier
    classifier.fit(X_train_scaled, y_train)

    # Prediction
    y_pred_test = classifier.predict(X_test_scaled)

    # Accuracy
    accuracy_test = metrics.accuracy_score(y_test, y_pred_test)
    accuracy.loc[iteration] = accuracy_test

```

```
# Display
print('Iteration', iteration, 'Accuracy =', '{:.8f}'.format(accuracy_test))

print('Mean accuracy', '{:.3f}'.format(accuracy.mean()))
```

```
Iteration 1 Accuracy = 0.93510324
Iteration 2 Accuracy = 0.92920354
Iteration 3 Accuracy = 0.95266272
Mean accuracy 0.939
```

4.3 Pipeline

```
[8]: from sklearn.pipeline import Pipeline
     from sklearn.model_selection import cross_val_score

     pipeline = Pipeline([('scaler', scaler), ('classifier', classifier)])

     accuracy = cross_val_score(pipeline, X, y, cv=cross_validation)
     print(accuracy)
     print('Mean test accuracy', '{:.3f}'.format(accuracy.mean()))
```

```
[0.93510324 0.92920354 0.95266272]
Mean test accuracy 0.939
```

```
[ ]:
```